

Latency: query/insert vs RAG context-prepend

