



The VLM can't identify



AgeDB cross-age, $N = 20$

The encoder can't localize



2-person scenes, $K = 2$