


CAPTIONING MEMORY

a group photo



"a woman, dark hair, mid-30s"

stored as text, searched later

months later



re-identified at 0.20



the description fits thousands

GROUNDED MEMORY (ours)

VLM grounds → encoder key

stored as one inline token

months later



recalled at 0.81



in-model, no round-trip

"remember her"