

Parametric Multimodal User Memory: Storing What Captions Cannot Carry

Bojie Li
Pine AI

Noah Shi
University of Washington

Abstract

A personalized agent needs a *user memory*: a persistent model of who its user is. Today that memory is almost always *text* — transcripts and captions, retrieved by similarity search. This serves the *captionable* half of a person (“my cat is named Bibi”) but drops the *perceptual* half that no caption can hold: how a voice sounds, how a face reads across age and lighting, how tired someone sounds today. We measure this loss: across five modalities, a captioner’s re-identification note recovers a perceptual identity at as little as 0.11 of a purpose-built encoder’s recall, collapsing toward chance where the signal is not nameable.

We instead **ground** perceptual memory in the model, decomposing recall into two subproblems. A vision-language model grounds the referent in context (*what* and *where*); a purpose-built encoder turns that region into an identity *key* (*who*); and the key is stored as one inline token, read by attention inside generation with no round-trip. Neither suffices alone — the VLM identifies cross-age faces at 0.54 where a face encoder reaches 0.81, and the encoder recognizes a chosen referent in a two-person scene at 0.05 — but grounded, the system reaches the correct-region oracle (0.96), and the same decomposition carries to multi-speaker audio and video. The recognition core is *training-free*, reproducing the encoder’s recall on any frozen model at $\mathcal{O}(1)$ registration cost. On **PerceptMem** (12 domains, 1,080 tasks) we further find perceptual identity capacity-limited (recall $\approx \min(1, k/M)$ of the encoder’s ceiling) and exact facts binding-limited: the first belongs in a parametric bank, the second in a text store. The two memories compose without interference — an agent with both remembers what its user said *and* what they were like.




Figure 1. A vision-language model grounds the referent the user means in a cluttered scene (*what/where*); a purpose-built encoder turns the grounded region into an identity key (*who*); and ATTMem stores the key on a frozen language model, reading it back as a single inline token. Neither component suffices alone — the VLM is a weak identity encoder, and a bare encoder cannot pick the referent out of a scene; grounded, they recover the correct-region oracle (Figure 2). The mechanism is detailed in Figure 4.

1 Introduction

When an AI agent “remembers” something about the person it serves, what does it keep? In almost every system today, *text*. Mem0 [Chhikara et al., 2025], MemoryLLM [Wang et al., 2024], MemGPT / Letta [Packer et al., 2023, Letta Team, 2024], M3-Agent [Long et al., 2025], and the wider conversational-memory line [Wu et al., 2025] turn what the agent saw or heard into transcripts (via ASR [Radford et al., 2023]) and captions (via BLIP-2 [Li et al., 2023] or LLaVA [Liu et al., 2023]), store the fragments in a sentence-encoder vector index [Reimers and Gurevych, 2019], and search them later. Personalized vision-language systems (MyVLM [Alaluf et al., 2024], Yo’LLaVA [Nguyen et al., 2024], MC-LLaVA [An et al., 2024], Online-PVLM [Bai et al., 2025a], RAP [Hao et al., 2025]) keep the same textual spine: a *named* concept (“Bibi”) is the handle, the perception an afterthought.

This is exactly right for the *captionable* half of a person — “my cat is named Bibi,” “I am vegetarian” — which passes through text intact. It is wrong for the other half: how a voice sounds, how a face reads across age and lighting, how a brushstroke marks a painter. None of this survives text — “a brown-haired man” does not tell two brown-haired men apart — because the loss is a property of the channel, not the wording; a re-identification note from a strong captioner recovers a perceptual identity at as little as 0.11 of a purpose-built encoder’s recall (Section 5). This *perceptual* half has no home today: it is flattened into a caption and lost, or handed to a recognizer that returns a bare label like “speaker 47” the model never sees. The missing piece is a memory that stores a perception *as a perception*.

Our proposal: grounded, parametric, multimodal user memory. We keep each perception in its native modality and split remembering it into three steps (Figure 1): a vision-language model *grounds* the referent the user means in context (*what* and *where*); a dedicated encoder (ArcFace [Deng et al., 2019], ECAPA-TDNN [Desplanques et al., 2020], CLIP [Radford et al., 2021]) turns the grounded region into an identity *key* (*who*); and ATTMem, a small per-modality bank on a frozen language model, *stores* the key and reads it back by attention at the output head as a single marker token. Recall is then a token the model conditions on — produced




Figure 2. Neither component alone; grounding fixes both blind spots. Two *separate* experiments, each isolating one failure. *Left (identity)*: on the same AgeDB cross-age faces ($N=20$), the VLM’s own vision tokens are a weak identity encoder (0.54) beside a purpose-built face encoder (0.81). *Right (localization)*: in two-person scenes ($K=2$), a bare encoder embedding the whole scene is near chance (0.05), while grounding the referent first recovers the correct-region oracle (0.96). The division of labour is a requirement, not a convenience: each component covers the other’s blind spot, and the detailed ablation follows in Figure 7.

inside generation with no caption and no retrieval round-trip, registered in one $\mathcal{O}(1)$ tensor append with no per-user training. ATTMEM borrows the attention-over-a-bank shape of k NN-LM [Khandelwal et al., 2020] and Memorizing Transformers [Wu et al., 2022], aimed at a new target: a user’s persistent perceptual identity.

Why grounding is necessary. A perceptual memory is not simply an encoder run over whatever the agent saw, because a raw perception mixes the identity worth remembering with the context around it. A friend photographed at a coffee shop and the same friend surfing share almost nothing pixel-for-pixel, so a whole-image embedding places them far apart — yet we recognise them at a glance by *looking at the face* and cropping the scene away, and only the cut-out face is invariant. Audio is the same: one voice on a phone call and on a podcast shares few words and different acoustics, while a listener locks onto the speaker. **Grounding is this act of cutting out the referent before encoding it**, and it forces a division of labour (Figure 2): the vision-language model finds *which* region to cut but is a weak identity encoder, while the specialist encoder is condition-invariant but cannot locate the referent in a cluttered scene. Grounded, each covers the other’s blind spot, recovering recognition neither reaches alone.

Figure 3 contrasts the two routes on one “remember her” episode.

Contributions.

- **Perceptual user memory must be grounded, not captioned.** We argue and measure that text cannot carry perceptual user content, and that recall must instead decompose into VLM grounding, a specialist encoder, and an in-model store (Sections 1–3).
- **A grounded memory that recovers oracle recognition in context.** The VLM grounds the referent, an encoder identifies it, and the identity is stored as an inline attention token; in cluttered scenes this reaches the correct-region oracle where whole-scene encoding is near chance — across vision, audio, and video (Section 5).




Figure 3. A worked example. “Remember her” in a group photo: a captioning memory stores a description that fits thousands and re-identifies her near chance later, while the grounded memory stores one inline token and recalls her cross-condition inside the model. Numbers are paired re-identification accuracies (Section 5).

- **A training-free, $\mathcal{O}(1)$ in-model recognition core.** Read by attention at the output head with *zero* gradient steps, it reproduces the encoder’s recall on ten frozen model families; insertion is $\mathcal{O}(1)$ and recall constant-time, far cheaper than keeping perceptions in context (Sections 3–5).
- **PerceptMem and a capacity-law router.** A benchmark of 12 perceptual domains and 1,080 tasks across five modalities, with text-only and chance baselines throughout, plus two capacity laws that route perceptual identity to the parametric bank and exact facts to a text store (Sections 4, 8).

Code. Our implementation and the PerceptMem benchmark are available at <https://github.com/19PINE-AI/multimodal-user-memory>.

2 Decomposing perceptual recall: grounding and identification

Why does the text route fail, not for lack of effort but structurally? Take the voice example in its hardest form: ten samples on file, a new one arriving cross-recording across different rooms and days, and the agent must name which of the ten it is. An agent has four options (Table 1). Three discard the very signal that tells one voice from another. The fourth keeps it — and exposes a second issue every route quietly conflates: recognizing a perception is really *two* problems, localizing the referent in context and encoding its identity, and no single component does both well.

The three that discard the signal. *Caption-and-search* [Chhikara et al., 2025, Wang et al., 2024] writes the voice down (“a male voice, mid-30s, slight Eastern-European accent”) and retrieves by similarity over the words. The description is true and useless. It fits thousands of speakers, so cosine over captions cannot separate ten users with even coin-flip reliability. *Recognize-and-label* [Long et al., 2025] calls a recognizer (ArcFace [Deng et al., 2019], ECAPA-TDNN [Desplanques et al., 2020]) and stores the label it returns (“speaker 47”). The model can recall that a label was assigned, but it never holds the perception and cannot compare today’s

Table 1. Four routes for remembering a perception, plus ours (citations in text). Only routes that keep the raw perception can discriminate across conditions. Embedding retrieval and ours both inherit the encoder’s recall exactly; ours additionally localizes the referent with a VLM and reads the identity back *inside* the model as a token the model conditions on, training-free and in $\mathcal{O}(1)$.

Route	What it stores	Keeps percept?	Cross-condition	Reg. cost
Caption-and-search	text description	no	chance	$\mathcal{O}(1)$
Recognize-and-label	a bare label	no	label only	$\mathcal{O}(1)$
Per-concept tuning	tuned token/adaptor	yes	strong	gradient / identity
Embedding retrieval	raw embedding	yes	encoder-limited	$\mathcal{O}(1)$
Parametric (ours)	embedding + LM value	yes	= encoder, in-model	$\mathcal{O}(1)$

voice against last week’s. *Per-concept tuning* [Nguyen et al., 2024, Alaluf et al., 2024] fine-tunes a token or adaptor per concept. It works, but each identity costs a second of gradient descent and yields a concept-specific artifact — the wrong cost profile for a companion registering dozens of characteristics per user, continuously.

The one that keeps it, and what it still needs. *Embedding retrieval* skips the words, indexes the raw voiceprint, and retrieves by cosine over the encoder’s vectors. It is the only text-free route that keeps the perception, and the strongest baseline we take seriously throughout. It succeeds exactly to the extent that the encoder is invariant across recording conditions, and fails to the extent that it is not. On a 2180-identity face pool it reaches 0.95 recall at ten identities — strong, but short of certainty, and eroding as the pool grows. Crucially, embedding retrieval presumes the perception has *already* been isolated: someone handed it a clean voiceprint or a tight face crop. In a real scene that isolation is itself a hard, context-dependent step, and it is where a bare encoder is helpless (a single embedding of a two-person photo recognizes the intended person at 0.05, Section 5). This is the second problem the four routes conflate: *localizing* the referent is not the same as *identifying* it, and the strongest text-free route only solves the second.

Our mechanism keeps this identity step and adds the two things it lacks: a VLM in front to resolve the referent, and a new home for the answer. Rather than return a neighbour from an external index, it stores the encoder key beside the model’s own marker-token vector and reads the pair back by attention inside the model (Section 3), so the recalled identity is a token the model conditions on. Recognition is unchanged — a sharp read approximates a hard argmax over the same encoder cosine — so what the memory adds is not a higher score but grounding, an in-model home, $\mathcal{O}(1)$ registration, and portability to any frozen model.

3 Grounding perceptions in the model’s own representation

Recall decomposes into three steps, and we take each in turn: *ground* the referent in context (the VLM), *identify* it with a dedicated encoder, and *store* the identity as a token read inside a frozen model. The first two are off-the-shelf components placed where they are strong; the third, ATTMEM, is the mechanism this paper contributes.

Grounding: resolving the referent in context. A perception rarely arrives pre-cropped. The user gestures at one person in a group, names the painting on the left, refers to the voice that just spoke; the system must turn that context-dependent reference into a region before anything

can be encoded. This is exactly what a vision-language model is good at and an encoder is not: the VLM grounds the referent against the surrounding scene and the user’s words, returning a region (a bounding box, a span of audio). We use the VLM’s native grounding for this and re-detect inside the returned region where a modality detector exists (RetinaFace [Deng et al., 2020] for faces, from the same InsightFace project as ArcFace), so box imprecision and un-aligned scenes are handled before identity is read. The necessity of this step is empirical: encoding the whole scene, with no grounding, recognizes the intended referent at 0.05 on faces and 0.11 on paintings, near chance, because a single embedding cannot represent a chosen referent among distractors (Section 5).

Identification: a dedicated encoder, not the model’s own tokens. The grounded region is read by a dedicated perceptual encoder into an L2-normalized identity *key*: ArcFace for faces, ECAPA-TDNN for speakers, CLIP for painting style, AST [Gong et al., 2021] for acoustic scenes, a wav2vec2 emotion encoder [Baevski et al., 2020] for tone. It is tempting to skip this and read identity from the VLM’s own context-conditioned representation — to have the model “output the embedding of the face part” directly. We tested that and it does not hold: a linear probe on a context-conditioned VLM hidden state recovers the referent at 0.25 (raw hidden state 0.14, against an oracle of 0.80), and the VLM’s native vision tokens recognize AgeDB cross-age faces at 0.54 where ArcFace reaches 0.81 at the same $N=20$ (Figures 2, 10); the gap persists end-to-end inside a real VLM (§6). The VLM is a strong localizer and a weak identity encoder, so the key must be a dedicated perceptual encoder, separate from the model’s own representation — which is also why an omni model’s encoder is best used as an external feature, not its in-context tokens.

Storage: an in-model attention memory. The identity key is stored and recalled by ATTMEM, a table of rows on a frozen language model (Figure 4). Row i holds a *key* k_i (the encoder embedding above) and a *value* v_i (the language model’s own embedding for a marker token assigned to that identity). Storing the model’s native vector as the value, rather than an arbitrary learned one, is deliberate: it makes the eventual logit boost for the right marker clean and self-reinforcing, rather than an accident of two unrelated vectors.

Reading the memory is one step of attention. When a localized perception arrives at generation time, the model forms a query q from its encoder key and, just before the output head, computes

$$w = \text{softmax}(\beta q^\top K), \quad r = w^\top V, \quad h' = h + g \cdot W_o r,$$

where K, V stack the bank’s keys and values, β is the attention sharpness, g a residual gain, and W_o an optional output projection — the identity in all our experiments, kept in the equation only to mark where a projection would sit. In words: the perception softly looks up the bank, pulls back a blend of the matching markers’ model-side vectors, and nudges the next-token prediction toward them. The blend is the entire read, and it adds *no trainable parameters*: β and g are two constants, and the values are the model’s own marker embeddings.

Design choices for a faithful read. Four details determine whether the read recovers the encoder’s nearest neighbour cleanly, and getting them right is what lets the mechanism work with no training. (i) *Do not* divide the attention logits by $\sqrt{D}$: with normalized keys, that




Figure 4. The grounded memory in detail. **Ground + identify (top):** the VLM resolves the referent to a box, a modality detector re-detects and aligns the region (RetinaFace for faces), and a frozen perceptual encoder (ArcFace / ECAPA / CLIP) maps it to an L2-normalized key $q=k \in \mathbb{R}^D$. **Store (bottom):** ATTMEM holds keys K ($N \times D$) and values V ($N \times H$, the model’s own marker-token embeddings); the attention read $w=\text{softmax}(\beta q^\top K)$, $r=w^\top V$, $h'=h+g W_o r$ adds a residual at the output head that biases the next token toward the matching marker. The read has no trainable parameters (β, g are constants); registration is a single $\mathcal{O}(1)$ append.

flattens every cosine difference into near-uniform attention. (ii) Set the sharpness β high, so the attention is a near-hard argmax over the bank rather than a diffuse average. (iii) Attach the hook at the output head, not at an intermediate layer, so the residual reaches the logits undiluted. (iv) Set the gain g large, so the retrieved marker dominates the hidden state and the read is exact; on models with untied output embeddings, whose output-embedding norms run smaller, g must be larger still (Appendix A, Table 7). With these four choices the read reproduces the encoder’s recall exactly (§5.4), and registering a user is a single row append with no training at all.

Registration versus fine-tuning. The cost of writing to memory is a decisive property (Figure 5). Adding an identity is a single `torch.cat` (about half a millisecond to add a thousand), and recall is constant-time regardless of bank size, because the lookup is a small matrix multiply dwarfed by the model’s own forward pass (~ 15 ms either way). The natural alternative, feeding the registered perceptions into the model as context, grows linearly per query and exhausts memory past the context window. Our mechanism is $52\times$ faster at a thousand identities, and is the only one of the two that still functions at ten thousand. A memory that a companion updates after every conversation must be this inexpensive to write.




Figure 5. Recall latency is flat in bank size (~ 15 ms, dominated by the frozen model’s own forward pass) and insertion is $\mathcal{O}(1)$. Feeding the same registered perceptions in as context grows linearly and exhausts the 32k context window (Qwen2.5-3B) at ten thousand identities. Per-concept fine-tuning (PATH A-style) costs gradient steps per identity. Log-log axes.

4 PerceptMem: a cross-condition perceptual-memory benchmark

To measure perceptual recall we built **PerceptMem**, five sub-modalities each chosen for one reason: the signal that discriminates is provably destroyed by captioning (Table 2). They span vision and audio, identity and style, the physical and the affective:

Table 2. The five PerceptMem sub-modalities. Each is included because no caption can carry its discriminating signal.

Sub-modality	Source / encoder	Why text cannot carry it
Face across age & lighting	LFW+AgeDB / ArcFace	“brown-haired man” fits thousands
Painter style	WikiArt / CLIP-mid	no caption separates early vs. late Monet
Speaker across recordings	LibriSpeech / ECAPA-TDNN	timbre is not transcribable
Acoustic scene	ESC-50 / AST	“traffic noise” fits every street
Tone of voice	RAVDESS / wav2vec2	“sounded tired” lacks a personal baseline

The sources are all standard: LFW [Huang et al., 2008] and AgeDB [Moschoglou et al., 2017] read with ArcFace [Deng et al., 2019], WikiArt [WikiArt, 2010] with mid-layer CLIP [Radford et al., 2021], LibriSpeech [Panayotov et al., 2015] with ECAPA-TDNN [Desplanques et al., 2020], ESC-50 [Piczak, 2015] with AST [Gong et al., 2021], and RAVDESS [Livingstone and Russo, 2018] with a wav2vec2 emotion encoder [Baevski et al., 2020]. The benchmark deliberately tests *memory*, not perception: every sub-modality has a fixed, off-the-shelf encoder that any method may use as black-box infrastructure, so no one wins by having a better encoder. The interface is just `register` and `recall`. Data is cross-session: a registered sample and its query come from different recordings. The identities used to train the memory never overlap with those it is evaluated on. For a memory of N identities we register one sample each, then issue cross-condition queries and ask the method to name the right one. Throughout the paper, *recall* means recall@1: the fraction of queries for which the top-scoring registered identity is the

correct one — for retrieval, the encoder’s nearest neighbour; for the in-model read, the argmax over the registered marker logits. For tone of voice, each registered “identity” is a (speaker, emotional-state) pair, so recalling it means matching a paralinguistic state across separate utterances, something a per-utterance caption cannot do without the user’s own history.

For statistical rigor we expand these five conceptual sub-modalities into a 12-domain, 1,080-task benchmark across multiple datasets and encoders — each domain a (dataset, encoder) pair, enumerated with its pool size and per-domain recall in Appendix Table 4 (Section 5, Figure 10) — and report a text-only baseline and a chance floor on every comparison. Throughout, our reference point is **embedding retrieval**: the same encoder, cosine nearest-neighbour over the registered keys. We call it the *encoder ceiling*, because it is the best one can do with the encoder’s own similarity, and our in-model read reproduces it rather than exceeding it.

5 Empirical evaluation

We organize the evaluation around the claims of the grounded memory. First, the half text drops is real and large: captions cannot carry perceptual identity (§5.1). Second, grounding recovers it: VLM grounding plus an encoder identity reaches the correct-region oracle in cluttered scenes (§5.2), and the same decomposition generalizes to audio and video (§5.3). Third, the recognition core is the encoder, read in-model, which is a feature, not a shortfall (§5.4). We then show that the in-model store enables composition with text memory (§5.5), ties the strongest pipeline baseline while beating in-context images on both accuracy and cost (§5.6), and preserves text behaviour exactly (§5.7); finally, we assemble the pieces into an end-to-end multi-session agent (§5.8).

5.1 Text captions cannot carry perceptual identity

The baseline a system would otherwise use is a text note: describe the perception in words, store the sentence, retrieve by text similarity. We test this on *all five modalities* and give it its best shot — an LLM writes a re-identification note (Qwen2.5-VL [Bai et al., 2025b] for images, Qwen2.5-Omni [Xu et al., 2025] for audio), a sentence encoder embeds it, and recognition is cosine nearest-neighbour over the notes — scored against the encoder on identical draws. The result is a clean router gradient (Figure 6): the note’s recall slides from nearly the encoder’s where the signal is *nameable* (an acoustic scene, “someone typing on a keyboard”) down toward chance where it is a perceptual *identity* (a voice, whose “high-pitched, with an accent” fits countless speakers). Text is a competent memory for what can be said and a failing one for what cannot — precisely the line the router draws between the text store and the parametric bank (§8).

5.2 Grounding recovers oracle recognition in cluttered scenes

In deployment a perception rarely arrives pre-cropped: the referent sits among other people or objects, and the system must resolve the one the user means before encoding it. An ablation isolates the value of each step (Figure 7). A *text-only* memory, in which the VLM describes the referent from the whole scene and captions are matched, recovers it at 0.12 on faces and 0.17 on paintings; a *store-only* memory that embeds the whole scene is lower still, near chance, because a single embedding cannot represent a chosen referent among distractors. *Grounding* the referent first, then identifying it with the encoder, recovers the correct-region oracle in both




Figure 6. Text-only caption-and-search versus the parametric encoder, paired on identical draws across all five modalities ($N=20$, 30 draws, 95% CI); the label above each pair is the text/encoder ratio. Ordered by how nameable the signal is. Text nearly matches the encoder on the nameable acoustic scene (93%) and collapses toward chance on perceptual identity (speaker 11%), tracing the router’s boundary between the text store and the parametric bank.

domains — 0.96 on faces and 0.36 on paintings, each equal to its own oracle; the low painting ceiling is CLIP’s, not grounding’s — with the VLM resolving the right region on essentially every trial. Retrieval coincides with the store throughout — both are the encoder’s nearest neighbour — so the gain is grounding, not the matcher.

The advantage is robust to scene clutter and VLM scale (Figure 8). As the number of referents per scene grows from two to six, grounding accuracy stays at 1.00 through $K=5$ and slips only marginally at $K=6$ (0.99), with grounded recall tracking the oracle (0.95–1.00) throughout while whole-scene recall stays at the floor; text-only, by contrast, decays as clutter grows, dropping from 0.12 at $K=2$ to 0.05 at $K=3$ as the caption must disambiguate more distractors. Grounding accuracy is already saturated at 7B and unchanged at 32B (1.00 either way, recall 0.960 vs 0.958), so the decomposition does not depend on a large grounder. Grounding is lossless and domain-general; what it recovers is bounded only by the encoder, modest for CLIP style and near-perfect for ArcFace identity.

5.3 Grounding generalizes to audio and video

The decomposition is not vision-specific: the same three steps apply when the referent is a *speaker* in a multi-speaker recording or a *person* in a video (Figure 9). For **pure audio**, we concatenate K real VoxCeleb [Nagrani et al., 2017] speakers into one clip and ask an audio-LLM (Qwen2.5-Omni) to ground the referenced speaker’s time span — the direct analog of the VLM grounding a bounding box — then identify that span with ECAPA. Grounded recall equals the oracle segment at $K=2$ (1.00) and $K=3$ (0.92, grounding 0.96) and degrades at $K=4$ (0.67, grounding 0.69) as temporal grounding of the fourth speaker becomes hard, while whole-clip encoding stays near chance and the encoder itself is always perfect. The result carries to real conversation: on *VoxConverse* [Chung et al., 2020], in-the-wild YouTube conversations with natural turn-taking, overlap, and noise, grounding is harder (0.75) but the grounded read (0.73) still nearly doubles whole-window encoding (0.38) toward the oracle (0.97), across 60 conversations, and the *AMI* meeting corpus [Carletta et al., 2005] behaves the same (grounded 0.75 vs whole-window 0.34, grounding 0.88). This mirrors the vision finding exactly: grounding, not identity, is the component that degrades with complexity.




Figure 7. Ablation of the grounded memory on cluttered scenes ($M=40$, two referents per scene). Capability increases left to right: *text-only*, where the VLM describes the referent from the whole scene and captions are matched; *store-only*, embedding the whole scene with no grounding; *grounded* (ground + identify + store, ours); and the correct-region *oracle*. Both text-only and store-only are near chance, while grounding recovers the oracle. Retrieval coincides with store throughout (both are the encoder’s nearest neighbour), so the gain is grounding, not the matcher. Faces: ArcFace, landmark alignment, 5 seeds; paintings: CLIP, no alignment, 3 seeds; error bars $\pm 95\%$ CI.




Figure 8. The grounded advantage is robust to scene clutter and grounder scale (faces, $M=40$). **Left:** as referents per scene grow from two to six, grounded recall (navy) stays at the correct-region oracle (light blue, coincident) while whole-scene recall stays at the floor. **Right:** grounding accuracy stays at 1.00 through $K=5$ and dips only to 0.99 at $K=6$, and is identical for Qwen2.5-VL-7B and -32B, so the decomposition does not require a large grounder.

For **video** (vision and audio together), we use real RAVDESS recordings, in which the same 24 actors are captured on video across eight emotions, giving a genuine (face, voice) pair per identity. We register each actor from one clip (an ArcFace face key and an ECAPA voice key) and recognise a held-out clip of a different emotion. Here face recognition is near-saturated (1.00) and cross-emotion voice recognition is hard (0.21), so fusing the two channels with equal weight does not beat the face alone — and we do not claim it does. The value of storing *both* is robustness to a corrupted channel, which a single-modality memory cannot provide (Figure 9, right). When the face is degraded to a distant, low-resolution view, face-only recall collapses to near chance (0.09) but the fused memory tracks the surviving voice (0.21); when the audio is corrupted by noise, voice-only drops (0.16) but the fused memory tracks the surviving face (0.83). Fusion is the only representation that never collapses across conditions, which is




Figure 9. Grounding generalizes beyond vision. **Left — pure audio:** an audio-LLM grounds the K -th speaker’s time span in a multi-speaker VoxCeleb clip and ECAPA identifies it ($M=30$, 3 seeds). Grounded recall (navy) tracks the oracle segment at $K=2-3$ and degrades at $K=4$ as temporal grounding gets hard; whole-clip encoding stays near the floor. **Right — vision+audio (RAVDRESS video):** each identity carries a face and a voice key. Storing both is robust to single-channel corruption: face-only collapses when the face is a distant low-resolution view, voice-only is weak and drops under noise, but the fused memory never falls below the surviving channel ($M=20$, 40 draws).

precisely the condition a deployed companion faces as people turn off-camera or speak over noise.

5.4 The in-model read reproduces the encoder’s recall

Once the referent is localized and encoded, recognition is the bank’s job, and the bank reproduces the encoder. With a sharp read the attention approximates a hard argmax over the same encoder cosine that embedding retrieval uses, so its recall matches the encoder’s, neither better nor worse. A paired protocol that scores both on identical registrations and queries with the target slot randomised confirms this at every bank size from $N=5$ to $N=1000$ ($|\Delta| \leq 0.001$ at each N , 20 draws; Table 5, Figure 17a): the memory ties retrieval on random banks and on adversarial banks of near-identical look-alikes, and a learned metric on the same features lands in the same place (Appendix A). The tie depends on the read being sharp: run with an under-sharpened read, the same sweep trails the cosine at $N \geq 300$ (Appendix A), a gap the constants of Section 3 eliminate.

Breadth and statistical rigor. At scale the recognition core holds across 12 domains and five modalities — 1,080 tasks ($N \in \{10, 20, 40\}$, 30 draws, 95% CIs), clearing chance in every one (Figure 10). Recall *tracks the encoder*: near-perfect where the encoder is strong (speaker, face), low where it is weak (style, tone), and, tellingly, the VLM’s own native vision encoder reads the same faces far below a purpose-built one — the cleanest statement that the encoder, not the memory, sets each ceiling (§6).

Training-free on any frozen model. Because the read is a similarity-weighted blend of marker tokens, nothing about the recall depends on the host language model, and it ports without training. Across ten families (Qwen2.5 1.5B/7B [Yang et al., 2024], Qwen3 4B/8B/14B [Yang et al., 2025], Phi-3.5 [Abdin et al., 2024], SmolLM2 [Ben Allal et al., 2025], DeepSeek-Llama-




Figure 10. The recognition core across 12 perceptual domains and 5 modalities (1,080 tasks; recall@1 at $N=20$, 30 draws, 95% CI), coloured by modality, against chance ($1/N$, black ticks) and the text-only caption baseline where measured (purple diamonds, per modality). Parametric recall tracks the encoder and clears chance everywhere, far exceeding text-only on perceptual identity (speaker, face, tone). The one inversion is the *nameable* acoustic scene, where a caption (“typing on a keyboard”) edges the weaker AST encoder — exactly the router boundary (§8): text suffices for nameable categories, the parametric bank for identity. The VLM’s own native vision encoder (orange) reads the same AgeDB faces far below a purpose-built one.

8B [DeepSeek-AI, 2025], Mistral-7B [Jiang et al., 2023], and a hybrid-Mamba Granite [IBM Granite Team, 2025]) the read reproduces the encoder’s recall (Appendix Figure 16), and a five-architecture by five-modality grid holds to within one point in every cell (Appendix Table 6). On tied-embedding models the match is within 0.003 in every cell (and exact at every gain we tested on the face task, Table 7); untied models match as closely with one larger read-strength constant; the hybrid-Mamba reader is within a point. This universality is a supporting property of the store, not the headline: it says the “store” step is free to host on whatever model the agent already runs.

Rejecting strangers and verifying identities. A user memory must do more than closed-set recall, in which the query is always one of the registered identities. It must reject a stranger it has never enrolled, rather than assign them a registered name, and decide whether two perceptions are the same person. Both fall out of the same read (cosine over the encoder keys), and we report them with standard operating-point metrics over 30 draws (Figure 11). On verification the memory separates same- from different-identity pairs with AUROC 0.99 on voice and 0.97–0.99 on face (equal-error rate 1.3% and 5–8%); on open-set identification it rejects never-enrolled strangers at 0.96–0.99 while still identifying the enrolled. As everywhere, these track the encoder: they are high where it is strong (voice, face) and low where it is weak (painting style, open-set AUROC 0.54). The point is not a new number but that the memory supports the open-set and verification operating modes a deployed companion requires, inheriting the encoder’s separability for both.




Figure 11. Beyond closed-set recall, the memory verifies identities and rejects strangers ($N=20$ enrolled, 30 draws). Per modality: verification AUROC (same vs different identity), open-set AUROC (enrolled vs never-enrolled probe), and the reject-stranger rate at the balanced threshold. All are high where the encoder is strong (voice, face) and fall where it is weak (painting style), the same dependence as recall.




Figure 12. One recall, three views (held-out faces, ten identities). **(a)** with a sharp temperature, bank attention is a near-pure diagonal (avg. 0.98), committing to the top encoder match. **(b)** the underlying encoder cosine has a softer diagonal (avg. 0.46) with off-diagonal mass; the memory inherits whatever errors live here. **(c)** the resulting marker logits keep the argmax identity on top. The attention faithfully reads the encoder’s nearest neighbour as a token; it does not sharpen the similarity.

5.5 Composition with in-context text memory

The reason to recall a perception as a token, rather than a label from an index, is that the token participates in the model’s reasoning. We register M faces with name markers, bind one fact per name in context (“Anna is a teacher”), then show a held-out cross-condition photo and ask for the fact. The whole chain runs in one forward pass — the face recalls the name through the bank, the name retrieves the fact from context — and end-to-end accuracy beats the face-withheld chance baseline by 4–10 $\times$, tracking recognition times the model’s in-context lookup (Figure 13). What bounds the chain as M grows is that in-context lookup, not the memory. Figure 12 opens up a single read: a sharp temperature concentrates the encoder’s speckled cosine into a near-one-hot diagonal, so the marker logits put the argmax on the same identity retrieval would pick, as a decisive token rather than a diffuse blend.




Figure 13. In-model composition: register M faces with name markers, bind one fact per name in context, then recall the fact from a held-out cross-condition photo, all in one forward pass. End-to-end accuracy (gold) tracks recognition (navy) times the model’s in-context lookup (blue) and beats the face-withheld chance baseline by 4–10 $\times$. 12 draws; bands are $\pm$ std.

5.6 Comparison with retrieval and in-context baselines

The natural baseline is retrieve-and-reprompt: cosine-retrieve the nearest registered face, recover its name, and re-prompt the model. Because both routes recognise with the same encoder cosine and reason with the same frozen model, they reach *identical accuracy* — on composition and on a multi-perception “how many distinct people are here?” task alike (Table 3). On cost the pipeline is in fact *cheaper* for pure recognition, since counting distinct identities needs only cosine and no model forward at all. Against this strong baseline the parametric memory is not a faster or more accurate recogniser, and we do not claim it is. Its advantage is architectural, and it is worth stating exactly when that matters. The pipeline is a two-pass design: an orchestrator must decide *before* generation that a recognition will be needed, query the index, and re-prompt with the answer. The in-model read makes no such decision — the bank is consulted at every perceptual position of the generation the agent is already running, so recognition composes with whatever the model is midway through producing (§5.5), interleaves across many perceptions in one stream, and ships as a single frozen-model artifact with no external index to host, synchronize, or secure. Where a deployment can always afford the second pass, the pipeline is a fine choice; the parametric memory is for the agent that cannot know in advance when a perception will need a name.

The comparison sharpens against the baseline a VLM actually reaches for: keeping the perception in context as image tokens. A face crop costs 64 vision tokens in Qwen2.5-VL at 224px and 256–576 at higher resolutions (a full image approaches a thousand), where ATTMEM reads the same perception as *one* marker token. That is a 64–576 $\times$ reduction in prefill tokens per perception, and because the bank lives outside the context the memory scales without bound while image-in-context exhausts a 128k window at about two thousand crops. On accuracy this baseline is also far weaker, since the VLM recognises with its general vision tokens (0.50 end-to-end, Appendix Table 11) rather than a face encoder (0.95). So the parametric memory ties the external-index pipeline and outperforms the in-context-image baseline on both accuracy and cost. That, together with the capacity-law router, is the contribution, not a benchmark number.

Table 3. In-model memory vs. a retrieve-and-reprompt pipeline (Qwen2.5-3B, 10 draws). Accuracy is identical everywhere, because both recognise with the same encoder cosine. The pipeline is cheaper for pure recognition (the distinct-count task needs no model forward); the in-model read folds recognition into the model’s forward instead of an external index.

Task	In-model acc	Pipeline acc	In-model lat.	Pipeline lat.
Compose (face→fact)	0.81	0.81	122 ms	98 ms
Count distinct, $K=2$	1.00	1.00	238 ms	0.09 ms
Count distinct, $K=3$	1.00	1.00	361 ms	0.09 ms
Count distinct, $K=4$	0.70	0.70	472 ms	0.09 ms

5.7 Preservation of text-only behaviour

The perceptual memory is meant to sit beside an ordinary text memory, not perturb it. The hook adds a residual only at perceptual positions, so with the bank installed and populated the model’s text-only behaviour is preserved: the top-1 next-token prediction matches the untouched model on every probe, the populated bank produces exactly the same logits as an empty one, and with the hook idle the logits are bitwise identical ($\max |\Delta \text{logit}|=0$; the residual 0.375 deviation on the wrapped forward comes from the host’s `inputs_embeds` numerics and is present with the bank empty or full alike). Register faces and voices together and the two banks stay independent without any training (cross-modal leakage 0.017 and 0.067). An agent can keep “my favourite restaurant is in Paris” in its text store and “how Bibi looks across lighting and age” in its perceptual bank, and lose nothing on either. The discrete-codebook design we tried before this one, which snapped each perception to one of K learned codes, never escaped ~ 0.07 recall past a few hundred identities for the same reason captioning fails — a categorical bottleneck discards the signal the encoder preserved; we keep it as a contrast in Appendix C.

5.8 An end-to-end multi-session agent

The pieces compose into a working agent. We build one that runs the full system — a frozen language model with the ATTMEM perceptual bank and a text store of facts — and enroll a growing population of users across sessions, each with a face (registered from one photo) and a bound fact. A returning user is perceived cross-condition (a different photo), and the agent must identify them, answer a question about them, and reject anyone it has never met. We evaluate as the enrolled population grows from ten to eighty users, over five seeds (Figure 14). The perceptual capabilities scale gracefully: identification holds at 0.86 and stranger rejection at 0.81 at eighty users, tracking the encoder as everywhere (the 0.96–0.99 rejection of §5.4 is measured at twenty enrolled; the operating point tightens as enrollment grows, exactly as the encoder’s separability does). Fact recall exposes the router directly. Routed through the text store — the perception recalls the name, the name keys an exact fact lookup — accuracy tracks identification (0.86 at eighty users). Attempted instead *in-model*, by binding all facts in the prompt and reading the answer in one pass, it collapses beyond a few dozen users (0.55 at twenty-five, 0.04 at fifty), because a 3B model cannot hold fifty name–fact bindings in context. This is the router thesis made concrete: perceptual identity belongs in the parametric bank, where it scales, and exact facts belong in a text store, where lookup stays exact. The production path — perceptual bank for *who*, text store for *what* — sustains end-to-end task success at 0.83




Figure 14. An end-to-end multi-session user-memory agent (Qwen2.5-3B, 5 seeds) as the enrolled population grows. Identifying a returning user (face→name), recalling their fact through the router (name→text store), and rejecting strangers all scale gracefully, tracking the encoder. Recalling the fact *in-model* instead, with all facts bound in context, collapses beyond a few dozen users — the host model’s in-context limit, not the memory’s — which is why the router keeps facts in a text store. End-to-end success holds at 0.83 for eighty users in one frozen model.

for eighty users inside a single frozen model, with no external index.

6 The encoder determines recall

Because the read reproduces the encoder’s recall, the memory’s quality *is* the encoder’s quality, and recall moves with encoder quality on the *same* identities: on AgeDB, ArcFace reaches 0.95 at ten identities and 0.75 at a hundred, while a VLM’s own vision tokens manage only 0.64 and 0.34 on those exact faces, and a learned metric on the same features does no better (Appendix A). This is why the “identify” step uses a dedicated encoder rather than the model’s native tokens — a gap that persists when the memory runs end-to-end inside a real VLM (§5.6, Appendix Table 11). The store, in contrast, is model-agnostic: the read carries no encoder-specific machinery, so it ports to any frozen model without training (§5.4). If a task needs better perceptual recall, the only remedy is a stronger encoder — not a different language model or a cleverer read.

7 Related work

Text-first agent memory. The conversational-memory line — Mem0 [Chhikara et al., 2025], MemGPT and its Letta successor [Packer et al., 2023, Letta Team, 2024], MemoryLLM [Wang et al., 2024], M3-Agent [Long et al., 2025], benchmarked by LongMemEval [Wu et al., 2025] — stores transcripts and captions in a text index and retrieves by similarity. These systems define the captionable half of user memory that ours composes with; none keeps a perception in its native modality, which is the half this paper adds (§5.1 measures what that omission costs).

Personalized vision-language models. MyVLM [Alaluf et al., 2024], Yo’LLaVA [Nguyen et al., 2024], MC-LLaVA [An et al., 2024], RAP [Hao et al., 2025], Online-PVLM [Bai et al., 2025a], and TAME [Hong et al., 2026] personalize a VLM to user-specific concepts, most by learning an artifact per concept — a classifier head, a soft token, a projection. Our setting

differs in cost profile and mechanism: registration is a training-free $\mathcal{O}(1)$ row append, and Appendix Table 10 finds that the learned per-concept artifacts do not buy accuracy over the encoder’s own cosine on cross-condition identity, while costing per-identity gradient descent.

Retrieval inside the model. ATTMEM borrows the shape of attention over an external bank from k NN-LM [Khandelwal et al., 2020], Memorizing Transformers [Wu et al., 2022], and RETRO [Borgeaud et al., 2022], which retrieve from token or chunk stores to improve language modelling. We retarget the mechanism: the keys are perceptual-encoder identities rather than hidden states of past text, the bank holds a user’s persistent percepts rather than a corpus, and the read lands at the output head of a frozen model with no training.

Perception-to-token bridges and speaker recognition. Flamingo’s resampler [Alayrac et al., 2022] and BLIP-2’s Q-Former [Li et al., 2023] learn to project perception into a language model’s token space; our read instead injects a single marker whose value is the model’s *own* embedding, which is what makes it training-free. The speaker-recognition and diarization literature [Nagrani et al., 2017, Desplanques et al., 2020, Park et al., 2022] solves who-spoke-when as a standalone pipeline returning labels; we consume the same encoders as grounded identity keys the model conditions on, rather than as an external verdict.

8 Discussion

The claim is structural, and the numbers illustrate it. The argument is not “our method scores higher.” It is that text is the wrong container for half of what an agent should remember about a person, that recognizing a perception decomposes into grounding it and identifying it, and that a frozen model fronted by a VLM and backed by an encoder bank solves both halves where text and a bare encoder each solve neither. The empirical results matter because they substantiate each piece: the text loss is real and modality-graded (Figure 6), the grounding step is necessary and lets the system recover the oracle (Figure 7), and the recognition core is the encoder’s recall given an in-model home, training-free and universal (Appendix Figure 16). The contribution is the grounded, in-model container, not a recall score; the score is the encoder’s.

This is additive infrastructure, not a replacement. The perceptual memory is designed to bolt onto the text memory that agents already have [Chhikara et al., 2025, Wang et al., 2024, Packer et al., 2023, Letta Team, 2024], and the non-regression result is what makes that safe: text recall is preserved exactly, and the perceptual bank activates only at perceptual moments. As long-term memory benchmarks like LongMemEval [Wu et al., 2025] grow to include perceptual content, the gap this paper argues for becomes directly measurable.

The captionable half is going parametric too. This paper is one half of a larger bet: a personalized agent should keep its memory *inside* the model, not in an external index. A companion line makes the same bet for the *captionable* half. *User as Engram* [Li, 2026c] writes per-user *facts* into a hash-keyed parametric memory, showing that the obvious alternative (a per-user LoRA adapter [Hu et al., 2022]) is *reasoning-negative*: it memorizes facts but degrades the reasoning that uses them, because a LoRA edit is global where a row insertion is local. *User as Code* [Li, 2026b] pushes the same half into executable typed state, and *Programmable KV Cache* [Li, 2026a] keeps it as editable, composable notes in the attention cache rather than in weights or an external index. Our perceptual counterpart shares that thesis with a different mechanism, continuous cross-attention over encoder banks rather than hash-keyed rows. That continuous channel, carrying a perception as a vector projected into the model’s own




Figure 15. Two capacity laws for memory compressed into k soft tokens read by a frozen language model, the mechanism behind the router. **Left:** perceptual identity is capacity-limited. Compressing M faces into k prototype slots and recognising a cross-condition query gives recall $\approx \min(1, k/M)$ times the encoder ceiling $C(M)$ (≈ 1 for faces), so a slot per identity recognises everyone and more slots hold more people (voices behave identically). **Right:** exact factual content is binding-limited. A latent holds one exact code (0.87) but multi-code retrieval collapses at two (0.06) and adding latent tokens does not rescue it. Recognition scales with the latent budget. Exact recall does not, so it belongs in a text store.

embedding space rather than as text, is the same bet *The Latent Bridge* [Li and Shi, 2026] makes for inter-model communication, where a learned latent link beats writing and re-reading text. The two halves compose into one memory that holds both what the user said and what they are like.

The router: what a latent memory can and cannot hold. The split between the two stores is not a convention. It follows from two capacity laws that we measured by compressing content into k soft tokens read by a frozen language model (Figure 15). Perceptual identity is capacity-limited and degrades gracefully. Compressing M registered identities into k prototype slots and recognising a cross-condition query keeps a k/M fraction of them distinguishable, on faces, voices, acoustic scenes, painting styles, and tones of voice alike (Appendix Table 8). One slot per identity recognises as many as the encoder can, and fewer slots lose identities in proportion. The general law is recall $\approx \min(1, k/M) \cdot C(M)$, where $C(M)$ is the encoder’s own recall at M registrations: near 1 for strong encoders (face, voice), lower for the weaker style and tone encoders, so the slot-compression factor $\min(1, k/M)$ is what is universal and the ceiling is the encoder’s. This is not an artifact of k -means: a compressor whose k slots we instead learn by gradient descent lands on the same curve, because k hard slots can keep at most k of M identities separable (a pigeonhole bound). Exact factual content behaves in the opposite way. A latent holds one exact code and reads it back (0.87 for a six-character code), but storing several and retrieving one by name collapses at once (0.06 at two codes, near zero beyond), and adding latent tokens does not help. Recognition is a capacity problem, where a latent yields one identity per slot; exact recall is a binding problem, where a latent cannot match a key to its value at any budget. That is why the captionable half belongs in a text store, which performs that lookup exactly, and the perceptual half in a parametric bank, which scales with its slots — the same boundary the text ablation drew empirically (Figure 6), now derived.

Privacy and responsible use. A perceptual user memory stores biometric templates — faceprints and voiceprints — and must be held to the standards that govern them. Enrollment should be consented and explicit (the “remember her” act of Figure 3 is an affirmative request,

not passive collection), the bank kept on-device or encrypted at rest, and the same mechanism must not be repurposed for identifying people who never asked to be remembered. Two properties of the design help. Because one identity is one row, deletion is as cheap as registration: dropping the row removes the identity entirely, with no fine-tuned weights to unlearn — a clean right-to-be-forgotten that per-concept tuning cannot offer. And because recall *is* the encoder’s, the memory inherits the encoder’s documented demographic accuracy skews (face recognition most of all); deployments should evaluate the chosen encoder on their user population and prefer verification-style operating points (Figure 11) with stranger rejection enabled, so an unenrolled person is refused rather than misassigned an enrolled name.

Limitations. (i) The memory cannot exceed the encoder. Its recall *is* the encoder’s — paired to within ± 0.001 from $N=5$ to $N=1000$ (Table 5, Figure 17a) — so at very large memories it declines exactly as the encoder’s cosine does (0.78 at a thousand faces), and no in-model trick recovers more. (ii) Our largest cross-condition pool holds 2,180 real identities, from which we measure recall to $N=1000$ registered; the constant-time behaviour is confirmed to ten thousand rows, but recall at that scale awaits data. (iii) The grounding step inherits the grounder’s referring-expression resolution, and it is the component that degrades with complexity: on well-separated visual composites it is near-lossless to six referents (grounding 1.00 through $K=5$), but a stress test that pastes jittered, rescaled faces onto photographic backgrounds drops grounding to 0.63–0.81 and end-to-end recall well below its oracle (Appendix Table 12), and audio grounding of the fourth speaker falls to 0.69; denser, occluded, or noisier scenes will stress it further, and detector and diarizer robustness — not the encoder — is the practical bottleneck there. (iv) Inside a VLM, native vision tokens are weak identity keys (0.50 vs. ArcFace’s 0.95 at $N=10$, Table 11), so an external, modality-specific encoder is required. (v) The one per-model constant (read-strength gain) is set by inspection; untied-embedding models need a larger value, which we have not automated. (vi) The closest prior systems (MyVLM [Alaluf et al., 2024], Online-PVLM [Bai et al., 2025a]) have released neither code nor checkpoints; we compare against charitable reimplementations of their core mechanisms (Appendix A, Table 10).

9 Conclusion

A photograph and a written description of a face are not two encodings of the same thing. The description has already thrown away what lets you recognize the person. Agent memory today keeps only the description; we have argued it must keep the photograph too, and that keeping it well means grounding it. A vision-language model resolves the referent the user means; a purpose-built encoder turns it into an identity that survives the conditions text cannot; and a frozen model stores that identity as a content-addressable row it reads inside its own forward pass. Each component does what the others cannot — the VLM cannot identify, the encoder cannot ground, and an external index cannot return the answer inside the model — and together they recover oracle-level recognition that none reaches alone. The recognition itself is the encoder’s, given a home: training-free, constant-time, and portable to any frozen model. The two memories are complementary, text for what the user *said* and parametric multimodal memory for what the user *is like*, and an agent with both strictly dominates one with only text. That is where a face stops being a sentence, and starts being a memory.

Acknowledgements

We thank Xiangyu Zhang from StepFun AI for the initial ideas that inspired this paper; he introduced the problem of multimodal user memory to the authors and pointed out the limitations of purely textual memory. This paper was produced using Pine Copilot’s voice-directed *whisper coding* workflow [Pine AI, 2026], in which the authors specify, discuss, and review the work by voice while a coding agent (Claude Code with Claude Opus 4.8) carries out the planning, coding, experiments, and paper writing. We thank BSQL Networking for hosting the NVIDIA RTX PRO 6000 GPU.

A Detailed results

This appendix collects the per-cell numbers behind Sections 5–6. Recall is top-1 with the argmax restricted to registered markers. Tables were measured on different face pools, splits, and draw protocols (single-draw sweeps versus the 20-draw paired protocol below), so raw-cosine columns differ slightly across tables (e.g. 0.767 vs. 0.776 at $N=1000$); comparisons are valid within a table, not across tables. Tables marked *single eval draw* (Tables 9 and 10) are directional comparisons on one shared draw; the paired multi-draw protocol below is the statistical reference.

The 12 PerceptMem domains. Table 4 enumerates the benchmark behind Figure 10: each domain is a (dataset, encoder) pair, evaluated at $N \in \{10, 20, 40\}$ with 30 draws each (90 tasks per domain, 1,080 total). The five conceptual sub-modalities of Table 2 each appear under at least one dataset–encoder pairing, several under two, plus the VLM’s native vision tokens as the deliberately weak twelfth key space.

Table 4. The 12 PerceptMem domains: dataset, identity key encoder, cross-condition identity pool, and parametric recall@1 at $N=20$ (mean $\pm 95\%$ CI over 30 draws; chance 0.05). Tone identities are (speaker, emotional-state) pairs. The Qwen2.5-VL native-token row is included as a key-space control, not a recommended configuration (§6).

Domain	Key encoder	Pool	Recall@1, $N=20$
Face / AgeDB (cross-age)	ArcFace	500	0.869 ± 0.021
Face / LFW	ArcFace	1680	0.932 ± 0.018
Face / LFW	AntelopeV2	901	0.953 ± 0.012
Face / LFW+AgeDB combined	ArcFace	2180	0.921 ± 0.019
Speaker / LibriSpeech	ECAPA-TDNN	58	0.992 ± 0.004
Speaker / VoxCeleb (in-wild)	ECAPA-TDNN	40	0.991 ± 0.005
Acoustic scene / ESC-50	AST	50	0.792 ± 0.023
Painting style / WikiArt	CLIP (mid-layer)	128	0.319 ± 0.024
Painting style / WikiArt	DINOv2 [Oquab et al., 2024]	50	0.236 ± 0.023
Vocal tone / RAVDESS	wav2vec2 (emotion)	168	0.496 ± 0.021
Face / AgeDB (control)	Qwen2.5-VL native tokens	567	0.542 ± 0.029
Painting style / WikiArt	contrastive head	26	0.245 ± 0.018

Paired evaluation protocol. A comparison between the memory and retrieval requires two controls: both methods score on *identical* registrations and queries, and each target identity

occupies a randomly chosen bank slot. The second control matters because a fixed target slot lets the constant marker token in that slot accrue a baseline output-logit, which inflates recall independently of the query. We score both methods on the same twenty draws with the target slot randomised, and report mean $\pm$ std. Table 5 gives the result, with random-bank face rows spanning $N=10$ to $N=1000$. The memory ties retrieval at every size and in every regime, since with a sharp read it computes the same argmax over encoder cosine — even on banks of the nineteen most-confusable look-alikes, where the cosine gaps are smallest (0.853 vs. 0.853, $\Delta=0.0\text{pp}$). Recall is not monotone in N (0.749 at $N=300$ vs. 0.776 at $N=1000$): a draw at $N=1000$ spans nearly the whole 1090-identity eval pool, whereas smaller N are random subsets whose average difficulty differs from the pool’s; the effect is identical in both columns, so it is a property of the encoder and the pool, not of the memory. The sharpness is load-bearing: an earlier trained variant whose learned temperature settled at inv-temp ≈ 20 trailed retrieval by eight points on the look-alike cell and by nine to eighteen points at $N \geq 300$, because a soft blend of marker values carries no per-identity signal beyond the cosine it dilutes. With the constants of Section 3 the blend is a faithful argmax in every cell we measure.

Table 5. Paired evaluation: the memory and retrieval scored on identical registrations and queries across 20 draws, with the target slot randomised. The memory ties retrieval everywhere — random banks through $N=1000$, style banks, and adversarial banks of the nineteen most-confusable look-alikes. Recall is the encoder’s throughout.

Regime	Cell	Retrieval	ATTMEM (mean $\pm$ std)	Δ
random	Face, $N=10$	0.948	0.948 ± 0.033	0.0pp
random	Face, $N=100$	0.792	0.792 ± 0.021	0.0pp
random	Face, $N=300$	0.749	0.749 ± 0.018	0.0pp
random	Face, $N=1000$	0.776	0.777 ± 0.008	0.0pp
random	Style, $N=5$	0.473	0.473 ± 0.104	0.0pp
random	Style, $N=10$	0.428	0.428 ± 0.076	0.0pp
adversarial	Face, $K=19$	0.853	0.853 ± 0.007	0.0pp

Universality grid. Table 6 runs the training-free read across five model architectures and all five modalities, scoring each cell against the encoder with the same paired protocol (20 draws). Every one of the twenty-five cells reproduces the encoder’s recall to within one point (worst case 0.010, most cells ≤ 0.003). Tied-embedding models use the default read-strength gain; untied and hybrid models use a single larger value (256), set once per model and held fixed across modalities. The read therefore inherits the encoder’s recall regardless of the host model’s family, embedding scheme, scale, or even whether it is a transformer.

Read-strength robustness. The one per-model constant, the residual gain, is set by inspection rather than search because exact recall holds across a wide range (Table 7, face $N=50$, paired). A tied-embedding model is exact at every gain from 16 to 1024; an untied one (Mistral) is below the encoder until a threshold near 128–256 and then exact and stable through 1024. Any value in the upper range works for either, so there is no tuning to get wrong.

Capacity law across modalities. Table 8 runs the perceptual capacity probe (compress M registered identities into k prototype slots, recognise a cross-condition query, 5 seeds) on all five

Table 6. Universality grid: maximum $|\Delta|$ between the training-free memory and the encoder over $N \in \{5, 10, 20\}$ (paired, 20 draws), for five architectures $\times$ five modalities. Every cell is within one point of the encoder. “gain” is the single per-model read-strength constant. Encoder recall@1 at $N=10$: Face 0.95, Speaker 0.99, Acoustic 0.86, Tone 0.52, Style 0.43.

Model	Emb.	gain	Face	Speaker	Acoustic	Tone	Style
Qwen2.5-7B	tied	64	0.001	0.000	0.002	0.003	0.001
Phi-3.5-mini	tied	64	0.001	0.000	0.000	0.002	0.002
DeepSeek-Llama-8B	untied	256	0.000	0.000	0.002	0.002	0.000
Mistral-7B	untied	256	0.000	0.000	0.003	0.010	0.000
Granite-4.0 (Mamba)	hybrid	256	0.003	0.000	0.003	0.001	0.000




Figure 16. The training-free read reproduces the encoder on every frozen model, the supporting property behind the “store” step. **Left:** each model’s training-free recall lands on the encoder ceiling (paired, 20 draws). **Right:** the gap from the encoder at $N=50$ is ≈ 0 across ten families (1.5B–14B, tied and untied embeddings, hybrid-Mamba); only an untied model at the default gain dips (red), fixed by a single larger read-strength constant. No gradient steps anywhere.

Table 7. Read-strength robustness: training-free recall@1 (face, $N=50$, encoder ceiling 0.821) as the residual gain sweeps 16–1024. Tied models are exact throughout; untied models are exact above a threshold and stay exact. The constant tolerates a wide range.

Model	16	32	64	128	256	512	1024
Qwen2.5-3B (tied)	0.821	0.821	0.821	0.821	0.821	0.821	0.821
Mistral-7B (untied)	0.022	0.060	0.307	0.768	0.821	0.821	0.821

modalities. The slot-compression factor $\min(1, k/M)$ fits every modality once normalised by the one-slot-per-identity ceiling $C(M)$: the normalised error is ≤ 0.012 on the strong-encoder modalities and ≤ 0.13 on the weaker style and tone encoders. The ceiling $C(M)$ itself is the encoder’s recall and varies accordingly, from near 1 on voice to 0.26–0.36 at $M=32$ on style and tone. Recognition is capacity-limited and graceful in every modality; only the ceiling moves.

Learned-metric baseline. A natural question is whether *any* re-encoding of the same features beats raw cosine, even if our read does not. We fit a similarity on the same encoder features and the same identities (the 50/50 split, seed 42) and score it with the identical protocol. Two closed-form metrics suffice: regularised LDA and within-class whitening. Table 9 shows LDA

Table 8. Perceptual capacity law per modality. $C(M)$ is the one-slot-per-identity ceiling (recall at $k=M$), i.e. the encoder’s recall at M . “fit” is the mean $|\text{recall}/C(M) - \min(1, k/M)|$ over the $k < M$ cells: the slot-compression factor is universal; the ceiling is the encoder’s.

Modality	$C(2)$	$C(8)$	$C(32)$	fit to $\min(1, k/M)$
Face (ArcFace)	0.98	0.95	0.92	0.002
Voice (ECAPA)	1.00	1.00	0.99	0.001
Acoustic (AST)	0.95	0.87	0.76	0.012
Style (CLIP)	0.76	0.46	0.26	0.070
Tone (wav2vec)	0.94	0.67	0.36	0.125

edging raw cosine at small memory and falling below it past a few hundred identities, while whitening never moves more than a point from raw cosine in either direction. So there is no meaningful discrimination left to extract from these features: the encoder’s cosine is already the best ruler, which is why the memory’s inheriting it exactly is the ceiling, not a shortfall. A contrastive linear head we also trained underperformed both closed-form metrics and is omitted as a weak instantiation.

Table 9. Learned-metric baseline, recall@1, same encoder, same split, all methods scored on the same eval draw. LDA edges raw cosine at small N and falls below it at scale; whitening stays within a point of raw cosine throughout. No re-encoding of the same features meaningfully beats the cosine. ATTMEM’s recall equals raw cosine (paired through $N=1000$, Table 5), so the raw-cosine column doubles as its recall.

Modality	N	Raw cosine (= ATTMEM)	LDA	Whitening
Face (ArcFace)	5	0.933	1.000	0.933
Face (ArcFace)	10	0.933	0.967	0.933
Face (ArcFace)	300	0.734	0.667	0.744
Face (ArcFace)	1000	0.767	0.724	0.776
Style (CLIP-mid)	5	0.400	0.200	0.600
Style (CLIP-mid)	10	0.400	0.333	0.367

Comparison with per-concept methods. The personalized-VLM line (MyVLM, Yo’LLaVA, Online-PVLM) registers a concept by *learning* something per concept (a classifier head, a soft token, a projection). Their code and checkpoints are unreleased, so we reimplement each method’s core mechanism charitably on the cross-condition face task (ArcFace keys, same split). Table 10 reports accuracy and registration cost. Two findings. First, a per-concept linear classifier (MyVLM’s mechanism) is indistinguishable from raw cosine at every memory size (within single-draw noise), because with one registration embedding per identity the trained head collapses to the embedding itself. It yields no accuracy gain over retrieval, and it costs per-identity gradient descent that grows to 4.5 s at a thousand identities, against our $\mathcal{O}(1)$ tensor append. Second, a learned projection trained with supervised contrastive loss (Online-PVLM’s mechanism) *underperforms* raw cosine by 15 to 28 points on this task. We read this as a caution that a learned re-encoding is not without cost: tuned on identity-disjoint data, it can distort the very cross-condition geometry it was meant to sharpen. No method here, ours included, beats raw cosine on accuracy at scale; the parametric memory’s case rests on

cost and in-model behaviour, not recall.

Table 10. Charitable reimplementations of per-concept methods on cross-condition face recall (ArcFace, recall@1, single eval draw). MyVLM’s per-concept classifier ties raw cosine but its registration cost grows with the memory; Online-PVLM’s learned projection underperforms raw cosine. ATTMEM’s recall equals raw cosine (paired through $N=1000$, Table 5) at $\mathcal{O}(1)$ registration cost, so the raw-cosine column doubles as its recall.

N	Raw cosine (= ATTMEM)	MyVLM	Online-PVLM	MyVLM reg. cost
5	0.933	0.933	0.800	0.009 s
10	1.000	0.967	0.750	0.019 s
50	0.767	0.753	0.675	0.178 s
100	0.780	0.770	—	0.528 s
300	0.728	0.725	—	1.39 s
1000	0.776	0.778	—	4.52 s

Scaling and the codebook comparison. Figure 17a traces recall against memory size on the face pool under the paired protocol (20 draws per N). The training-free read is coincident with the encoder ceiling at every size — 0.948 vs. 0.948 at $N=10$, 0.777 ± 0.008 vs. 0.776 ± 0.008 at $N=1000$, $|\Delta| \leq 0.001$ throughout (Table 5) — so the memory declines exactly as the encoder’s cosine does, gracefully, from near-perfect at ten identities to 0.78 at a thousand. The coincidence requires the sharp read of Section 3: with a soft read (learned inv-temp ≈ 20) the same sweep trails the cosine by nine to eighteen points at $N \geq 300$, as recorded in the paired-protocol paragraph above. The discrete-codebook predecessor (PATH A), by contrast, flatlines at ~ 0.07 regardless of codebook size or budget. Figure 17b makes the same point across all five sub-modalities: replacing a discrete codebook with continuous attention lifts recall 2–10 $\times$, because a categorical bottleneck discards the signal the encoder preserved. This is the one place a design choice in the read matters; everything else inherits the encoder.




Figure 17. Left: under the paired protocol (20 draws per N) the training-free read is coincident with the encoder ceiling from $N=5$ to $N=1000$ ($|\Delta| \leq 0.001$; Table 5), both declining gracefully as the pool grows, while the codebook predecessor never gets off the floor. Right: continuous attention beats the discrete codebook 2–10 $\times$ on every sub-modality.

Vision-language model, key/value orthogonality. Table 11 gives the numbers behind the structural rule of Section 6. Each memory row is compared to retrieval *in its own key space*, under the training-free sharp read. An external ArcFace key (orthogonal to the model’s hidden space) reproduces the encoder’s recall exactly at every size — 0.948/0.792/0.776 at $N=10/100/1000$, the same values as on the text-LM hosts, as the universality result predicts. The VLM’s native vision tokens (already in hidden space) also tie their own retrieval — the read is faithful to whatever key it is given — but that ceiling sits far below ArcFace’s throughout, so they are the wrong key regardless of method: the empirical basis for reserving the VLM for localization and an external encoder for identity.

Table 11. Parametric memory inside Qwen2.5-VL-3B: same frozen model, only the key encoder changes; training-free sharp read throughout. ArcFace rows: paired protocol, 20 draws. Native-token rows: single draw on AgeDB with this harness’s own key pooling, so they differ from the AgeDB sweep of Section 6 (0.64/0.34); both agree that native tokens sit far below ArcFace. The read ties retrieval in both key spaces; only the key’s own ceiling differs. An earlier trained variant of the ArcFace row (12k steps, learned soft temperature) trailed retrieval at scale (0.494 at $N=1000$) — the same under-sharpening artifact recorded in the paired-protocol paragraph above.

Configuration	$N=10$	$N=100$	$N=1000$
Retrieval over ArcFace embeddings (encoder ceiling)	0.948	0.792	0.776
ATTMEM + external ArcFace keys (orthogonal to hidden space)	0.948	0.792	0.776
Retrieval over native vision tokens	0.50	0.43	—
ATTMEM + native vision tokens (co-located with hidden space)	0.50	0.42	—

Realistic-scene stress test. The composites of Section 5.2 place each face in its own well-separated cell, which is where grounding is near-lossless. Table 12 stresses the grounding step instead: faces are pasted onto a dimmed photographic background at jittered positions and varying scales ($0.7\text{--}1.05\times$), in a one-row or two-dimensional grid layout, and the referent is named by reading order (Qwen2.5-VL-7B, $M=40$, $K=3$, 3 seeds per layout). Grounding accuracy drops from 1.00 on clean composites to 0.81 (row) and 0.63 (grid), and end-to-end recall falls well below its oracle because the face detector also degrades on small, rescaled faces (hit-rate 0.09–0.43, forcing un-aligned fallback crops). This is the basis for limitation (iii): in cluttered photographic scenes the grounder and detector, not the encoder, are the practical bottleneck.

Table 12. Realistic-scene stress test: faces composited onto photographic backgrounds with positional jitter and scale variation (Qwen2.5-VL-7B, $M=40$, $K=3$; mean over 3 seeds). The correct-region oracle stays high, while VLM grounding and detector alignment degrade — the failure mode is localization, not identity.

Layout	Oracle region	Grounded (ours)	Whole-scene	Grounding acc.
Row (jittered)	0.83	0.42	0.03	0.81
Grid (jittered)	0.83	0.27	0.03	0.63

B Mechanics of a single recall

It helps to watch one recall end to end. Suppose the agent has met ten people and wants to register an eleventh from a single photo. Registration is one row appended to the face bank (Listing 1). The encoder turns the photo into a 512-dimensional ArcFace embedding, the *key*. The model’s own embedding for a freshly assigned marker token becomes the *value*, a vector of the model’s hidden width (2048 for Qwen2.5-3B). There is no gradient step and no fine-tuning. The row is the memory.

```
# Register identity #11 from a single photo -- no training:
k = l2_normalize(arcface(photo))      # key:   R^512  (encoder space)
m = assign_marker_token()             # e.g. "<id_11>"
v = model.input_embedding[m]          # value: R^2048 (model's own space)
bank.K = torch.cat([bank.K, k[None]]) # 0(1) append
bank.V = torch.cat([bank.V, v[None]])
```

Listing 1. One registered identity is one row. The key is the encoder’s view of the perception; the value is the model’s own vector for the marker token that names it. Adding the row is a single tensor append.

At recall, a new photo of one of the eleven arrives under different lighting and a year older. Its ArcFace embedding becomes the query q . Just before its output head, the model scores q against all eleven keys, softmaxes into attention weights w , and pulls back the weighted blend of values $r = w^\top V$. Because each value is the model’s own vector for a marker token, adding $g W_0 r$ to the hidden state lands as a clean boost on the right marker’s logit. The model’s next token *is* the remembered identity. The whole step is a few matrix multiplies dwarfed by the model’s forward pass.

The 0.98-vs-0.46 diagonal gap of Section 5 (Figure 12) is measured on exactly this probe: a held-out ten-identity bank, comparing the sharp attention weights against the raw encoder cosine they are computed from. The attention concentrates the cosine onto its top match, which is why the read recovers the encoder’s nearest neighbour as a single decisive marker rather than a diffuse blend.

C A discrete-codebook alternative

Before continuous attention we spent sixteen development cycles on PATH A: a discrete codebook that snaps each perception to one of K learned codes and remembers the code. It is the natural “compress the perception into a symbol” design, and it is a learned cousin of captioning, which is why its failure is instructive rather than embarrassing.

No setting broke it. Across codebook sizes $K \in \{128, 256, 512, 1024\}$, even after 100,000 steps of continual pretraining on the expanded face pool, top-1 recall at 300 identities stays pinned between 0.057 and 0.070 while embedding retrieval over the same encoder reaches 0.73. That is a roughly $10\times$ gap that does not move with K (Table 13). Swapping ArcFace R50 for AntelopeV2 R100 trained on 360K identities did not help either.

The reason is a vice the codebook cannot escape, and the diagnostics name it exactly. Two pressures pull in opposite directions as K grows. A *small* codebook packs many people into each cell: at $K=16$, distinct identities collide in the same code 8.7% of the time, and once two people share a code they are indistinguishable forever after. A *large* codebook fixes that, with

Table 13. PATH A after 100K-step continual pretraining on the 2180-identity face pool, at 300 registered identities. Recall is pinned near 0.07 regardless of codebook size, even when the gate routes to the correct code about half the time.

Codebook size K	128	256	512	1024
PATH A recall@1	0.057	0.064	0.070	0.070
gate routes to correct code	0.43	0.51	0.42	0.54
embedding retrieval (same encoder)	0.73			

inter-identity collisions falling to 2.4% at $K=128$, but it shatters each identity across cells. The rate at which two cross-condition photos of the *same* person land in the same code drops from 0.33 to 0.20, so the query no longer routes to where the registration was stored. Net recall is squeezed from both sides and never escapes ~ 0.07 . Even when the gate does route a query to the right code (about half the time, Table 13), every other identity sharing that cell is an equally good answer, so the final pick is near-random.

This is the same information loss as captioning, learned instead of written: any categorical bottleneck, whether a word or a code, throws away the continuous signal the encoder worked to preserve, and no amount of compute downstream can recover it. Continuous attention over the raw embedding never quantizes, and never pays this tax. PATH A is in the paper because its failure is the cleanest possible argument for the design that replaced it.

References

- Marah Abdin, Sam Ade Jacobs, Ammar Ahmad Awan, et al. Phi-3 technical report: A highly capable language model locally on your phone. *arXiv preprint arXiv:2404.14219*, 2024.
- Yuval Alaluf, Elad Richardson, Sergey Tulyakov, Kfir Aberman, and Daniel Cohen-Or. MyVLM: Personalizing VLMs for user-specific queries. In *European Conference on Computer Vision (ECCV)*, 2024.
- Jean-Baptiste Alayrac, Jeff Donahue, Pauline Luc, Antoine Miech, Iain Barr, Yana Hasson, Karel Lenc, et al. Flamingo: A visual language model for few-shot learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- Ruichuan An, Sihan Yang, Renrui Zhang, Ming Lu, Tianyi Jiang, Kai Zeng, Yulin Luo, Jiajun Cao, Hao Liang, Ying Chen, Qi She, Shanghang Zhang, and Wentao Zhang. MC-LLaVA: Multi-concept personalized vision-language model. *arXiv preprint arXiv:2411.11706*, 2024.
- Alexei Baevski, Henry Zhou, Abdelrahman Mohamed, and Michael Auli. wav2vec 2.0: A framework for self-supervised learning of speech representations. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- Huiyu Bai, Runze Wang, Zhuoyun Du, Yiyang Zhao, Fengji Zhang, Haoyu Chen, Xiaoyong Zhu, Bo Zheng, and Xuejiao Zhao. Online-PVLM: Advancing personalized VLMs with online concept learning. *arXiv preprint arXiv:2511.20056*, 2025a.
- Shuai Bai, Keqin Chen, Xuejing Liu, et al. Qwen2.5-VL technical report. *arXiv preprint arXiv:2502.13923*, 2025b.

- Loubna Ben Allal, Anton Lozhkov, Elie Bakouch, et al. SmolLM2: When smol goes big — data-centric training of a small language model. *arXiv preprint arXiv:2502.02737*, 2025.
- Sebastian Borgeaud, Arthur Mensch, Jordan Hoffmann, Trevor Cai, Eliza Rutherford, Katie Millican, George van den Driessche, et al. Improving language models by retrieving from trillions of tokens. In *International Conference on Machine Learning (ICML)*, 2022.
- Jean Carletta, Simone Ashby, Sebastien Bourban, Mike Flynn, Mael Guillemot, Thomas Hain, Jaroslav Kadlec, et al. The AMI meeting corpus: A pre-announcement. In *International Workshop on Machine Learning for Multimodal Interaction (MLMI)*, 2005.
- Prateek Chhikara, Dev Khant, Saket Aryan, Taranjeet Singh, and Deshraj Yadav. Mem0: Building production-ready AI agents with scalable long-term memory. *arXiv preprint arXiv:2504.19413*, 2025.
- Joon Son Chung, Jaesung Huh, Arsha Nagrani, Triantafyllos Afouras, and Andrew Zisserman. Spot the conversation: Speaker diarisation in the wild. In *INTERSPEECH*, 2020.
- DeepSeek-AI. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.
- Jiankang Deng, Jia Guo, Niannan Xue, and Stefanos Zafeiriou. ArcFace: Additive angular margin loss for deep face recognition. In *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019.
- Jiankang Deng, Jia Guo, Evangelos Ververas, Irene Kotsia, and Stefanos Zafeiriou. RetinaFace: Single-shot multi-level face localisation in the wild. In *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2020.
- Brecht Desplanques, Jenthe Thienpondt, and Kris Demuynck. ECAPA-TDNN: Emphasized channel attention, propagation and aggregation in TDNN based speaker verification. In *INTERSPEECH*, 2020.
- Yuan Gong, Yu-An Chung, and James Glass. AST: Audio spectrogram transformer. In *INTERSPEECH*, 2021.
- Haoran Hao, Jiaming Han, Changsheng Li, Yu-Feng Li, and Xiangyu Yue. RAP: Retrieval-augmented personalization for multimodal large language models. In *Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025.
- Rongpei Hong, Jian Lang, Ting Zhong, Yong Wang, and Fan Zhou. TAMEing long contexts in personalization: Towards training-free and state-aware MLLM personalized assistant. In *ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD)*, 2026. *arXiv:2512.21616*.
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations (ICLR)*, 2022.

- Gary B Huang, Marwan Mattar, Tamara Berg, and Eric Learned-Miller. Labeled faces in the wild: A database for studying face recognition in unconstrained environments. In *Workshop on Faces in Real-Life Images*, 2008.
- IBM Granite Team. Granite 4.0: Hyper-efficient, high-performance hybrid models for enterprise. IBM announcement, 2025. URL <https://www.ibm.com/new/announcements/ibm-granite-4-0-hyper-efficient-high-performance-hybrid-models>.
- Albert Q Jiang, Alexandre Sablayrolles, Arthur Mensch, et al. Mistral 7B. *arXiv preprint arXiv:2310.06825*, 2023.
- Urvashi Khandelwal, Omer Levy, Dan Jurafsky, Luke Zettlemoyer, and Mike Lewis. Generalization through memorization: Nearest neighbor language models. In *International Conference on Learning Representations (ICLR)*, 2020.
- Letta Team. Letta: Stateful LLM agents with long-term memory. Software framework, 2024. URL <https://github.com/letta-ai/letta>.
- Bojie Li. Models take notes at prefill: KV cache can be editable and composable. *arXiv preprint arXiv:2606.17107*, 2026a.
- Bojie Li. User as code: Executable memory for personalized agents. *arXiv preprint arXiv:2606.16707*, 2026b.
- Bojie Li. User as engram: Internalizing per-user memory as local parametric edits. *arXiv preprint arXiv:2606.19172*, 2026c.
- Bojie Li and Noah Shi. The latent bridge: A continuous slow-fast channel for real-time game agents. *arXiv preprint arXiv:2606.24470*, 2026.
- Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. BLIP-2: Bootstrapping language-image pre-training with frozen image encoders and large language models. In *International Conference on Machine Learning (ICML)*, 2023.
- Haotian Liu, Chunyuan Li, Qingyang Wu, and Yong Jae Lee. Visual instruction tuning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Steven R Livingstone and Frank A Russo. The Ryerson audio-visual database of emotional speech and song (RAVDESS). *PLoS ONE*, 13(5):e0196391, 2018.
- Lin Long, Yichen He, Wentao Ye, Yiyuan Pan, Yuan Lin, Hang Li, Junbo Zhao, and Wei Li. Seeing, listening, remembering, and reasoning: A multimodal agent with long-term memory. *arXiv preprint arXiv:2508.09736*, 2025.
- Stylianos Moschoglou, Athanasios Papaioannou, Christos Sagonas, Jiankang Deng, Irene Kotsia, and Stefanos Zafeiriou. AgeDB: The first manually collected, in-the-wild age database. In *CVPR Workshops*, 2017.
- Arsha Nagrani, Joon Son Chung, and Andrew Zisserman. VoxCeleb: A large-scale speaker identification dataset. In *INTERSPEECH*, 2017.

- Thao Nguyen, Haotian Liu, Yuheng Li, Mu Cai, Utkarsh Ojha, and Yong Jae Lee. Yo'LLaVA: Your personalized language and vision assistant. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- Maxime Oquab, Timothée Darcet, Théo Moutakanni, Huy Vo, Marc Szafraniec, Vasil Khalidov, Pierre Fernandez, et al. DINOv2: Learning robust visual features without supervision. *Transactions on Machine Learning Research (TMLR)*, 2024.
- Charles Packer, Sarah Wooders, Kevin Lin, Vivian Fang, Shishir G Patil, Ion Stoica, and Joseph E Gonzalez. MemGPT: Towards LLMs as operating systems. *arXiv preprint arXiv:2310.08560*, 2023.
- Vassil Panayotov, Guoguo Chen, Daniel Povey, and Sanjeev Khudanpur. LibriSpeech: An ASR corpus based on public domain audio books. In *IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*, 2015.
- Tae Jin Park, Naoyuki Kanda, Dimitrios Dimitriadis, Kyu J Han, Shinji Watanabe, and Shrikanth Narayanan. A review of speaker diarization: Recent advances with deep learning. *Computer Speech & Language*, 72, 2022.
- Karol J Piczak. ESC: Dataset for environmental sound classification. In *ACM Multimedia*, 2015.
- Pine AI. Pine AI: The most natural human-computer interface is your voice. Blog post, 2026. URL <https://www.19pine.ai/blog/pine-ai-the-most-natural-human-computer-interface-is-your-voice>. Accessed 2026-06-28.
- Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual models from natural language supervision. In *International Conference on Machine Learning (ICML)*, 2021.
- Alec Radford, Jong Wook Kim, Tao Xu, Greg Brockman, Christine McLeavey, and Ilya Sutskever. Robust speech recognition via large-scale weak supervision. In *International Conference on Machine Learning (ICML)*, 2023.
- Nils Reimers and Iryna Gurevych. Sentence-BERT: Sentence embeddings using Siamese BERT-Networks. In *Empirical Methods in Natural Language Processing (EMNLP)*, 2019.
- Yu Wang, Yifan Gao, Xiusi Chen, Haoming Jiang, Shiyang Li, Jingfeng Yang, Qingyu Yin, Zheng Li, Xian Li, Bing Yin, Jingbo Shang, and Julian McAuley. MemoryLLM: Towards self-updatable large language models. In *International Conference on Machine Learning (ICML)*, 2024.
- WikiArt. WikiArt: Visual art encyclopedia. <https://www.wikiart.org>, 2010.
- Di Wu, Hongwei Wang, Wenhao Yu, Yuwei Zhang, Kai-Wei Chang, and Dong Yu. Long-MemEval: Benchmarking chat assistants on long-term interactive memory. In *International Conference on Learning Representations (ICLR)*, 2025.

Yuhuai Wu, Markus N Rabe, DeLesley Hutchins, and Christian Szegedy. Memorizing transformers. In *International Conference on Learning Representations (ICLR)*, 2022.

Jin Xu, Zhifang Guo, Jinzheng He, et al. Qwen2.5-Omni technical report. *arXiv preprint arXiv:2503.20215*, 2025.

An Yang, Baosong Yang, Beichen Zhang, et al. Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*, 2024.

An Yang, Anfeng Li, Baosong Yang, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.